

01 The Opportunity

Who has this problem, what it costs them today, and how you would know if you fixed it.

A SUBMISSION DETAILS		Structured · Eligibility · pass/fail	
Complete every field. Incomplete submissions are screened out before scoring.			
FIELD	YOUR ANSWER	FIELD	YOUR ANSWER
Team name	SecureVoice AI	Contact email	otemaach@gmail.com
Track (1 or 2)	1 - Banking & Insurance	Based in	Kampala, Uganda
Use case (1-8)	1 - Real-time fraud intervention	Stage	Working prototype
Languages covered	EN, AR, HI, UR, FR, SW (6)	Prior ElevenLabs use (Y/N)	Y
Team size / based in	1 (solo founder) - Kampala, Uganda	Website or repo	github.com/BROCKUGANDA/SecureVoice-AI

B THE IDEA IN ONE LINE	25 words max · Problem fit · 25%
Complete the sentence: "An agent that ____ for ____, so that ____." No adjectives. An agent that verifies a flagged card and freezes it, in the customer's language, for bank fraud teams, so that prevented loss replaces write-offs.	

C WHAT BREAKS TODAY	120 words max · Problem fit · 25%
The current workflow and the exact point it fails. Name the channel, the delay and who is left waiting. Do not describe your solution here. A UAE retail bank's fraud engine flags a card transaction at 2 AM and the alert queues for the morning shift. An analyst calls at 9 AM - seven hours later. By then the fraudster has already called posing as the bank and moved 'safe' funds. The card is used again at 4 AM for AED 12,000; total loss AED 14,500, which the bank eats under CBUAE unauthorized-transaction rules. The customer loses trust and switches banks. The 60-second SLA fraud teams promise on paper does not exist in practice: the call channel is human, asynchronous and English-only.	

D TODAY'S BASELINE	Numbers only · Problem fit · 25%	
The measurable status quo. These exact numbers must reappear in box M — they will be cross-checked.		
WHAT YOU MEASURED	VALUE TODAY	WHERE THE NUMBER COMES FROM
Time from fraud signal to first customer contact	38 min	Fraud-ops SLA reports from the three box F banks
Languages live on fraud calls today	2 (EN/AR)	The same interviews: Arabic and English queues only
Unauthorized-transaction write-off rate	1.0%	Card-loss ratios in UAE bank annual reports; box F CRO

E WHO BUYS THIS	60 words max · Commercial · 20%
The institution, the job title that signs, and which budget line it comes out of. The Group Head of Fraud at a UAE retail bank signs, paid from the fraud-loss provision that already funds chargeback write-offs: AED 15,000 a month on voice intervention against AED 200,000 a month of write-offs. The CISO co-owns it - the platform sits inside the bank's risk perimeter and must clear the bank's own security review before pilot.	

02 The Evidence

Where your understanding came from, and how the work actually runs today.

F WHO YOU SPOKE TO				Structured · Problem fit · 25%		
Real conversations with people inside the institution type you are building for. We may contact them. Desk research alone scores in the bottom band.						
NAME AND ROLE	ORGANISATION TYPE	DATE	THE ONE THING THEY SAID THAT CHANGED YOUR IDEA			
Head of Fraud Operations	UAE retail bank (top 5)	Sep 2025	"We have the risk score, not the voice. The 38-minute delay is the call centre, not the algorithm."			
VP, Cards Risk	UAE Islamic bank	Sep 2025	"If the agent asks for a PIN even once the channel dies - our customers hang up on credential requests."			
Chief Risk Officer	UAE digital bank	Oct 2025	"Prove the call happened, prove what was said, prove no credentials were asked. CBUAE will ask all three."			

G WHAT YOU GOT WRONG				50 words max · Problem fit · 25%		
An assumption you held going in that turned out to be false, and what you changed. If nothing changed, you have not spoken to enough people. We assumed detection was the bottleneck. It is not - the call is. We assumed banks wanted a fully autonomous freeze agent; every bank we spoke to requires a human specialist before any irreversible action. So the agent executes only a pre-approved temporary hold, then hands off.						

H THE WORKFLOW TODAY				Diagram · Problem fit · 25%		
Draw the current process end to end, before your agent exists. Write into the lanes below — or replace this page with your own diagram at the same level of detail.						
	STAGE 1	STAGE 2	STAGE 3	STAGE 4	STAGE 5	STAGE 6
Customer or applicant	Card used at 2 AM	Asleep, unaware	Wakes to SMS at 9 AM	Bank calls, in English	Confirms or denies	Chases the chargeback alone
Front-line staff	-	-	Analyst picks up queue	Calls, asks questions	Freezes if confirmed	Escalates to specialist
Back office or approver	Risk engine flags	Alert queued overnight	Manual review	-	Freeze entered in core	Chargeback filed if loss
Systems touched	Core banking, fraud engine	Case management queue	Analyst workstation, CRM	Telephony (PSTN)	Core banking (freeze)	Chargeback system
Elapsed time	T+0	T+0 to T+7h (waits here)	T+7h	T+7h to T+9h (fails here)	T+9h	T+14 to T+45 days total

Your diagram must show:

☒ Every handoff between people or teams	☒ Each system by its actual name
☒ Where the process waits, and for how long	☒ The step that fails most often
☒ Where the customer has to chase	☒ Total elapsed time, start to finish

03 The Agent

How the conversation runs, what you build it with, and where the guardrails sit.

I THE CALL FLOW		15 words per step · Agent design · 25%	
One call, start to finish. Step 1 must be the opening disclosure. Mark the human handover with (H). Leave unused steps blank.			
STEP	WHAT HAPPENS		
1	Opening disclosure: recorded call, the bank's AI assistant identifies itself and states why		
2	Agent states the flagged transaction - merchant, amount, date - and asks: is this yours?		
3	'Not mine' - agent executes the pre-approved temporary card freeze and says so		
4	'It's mine' - agent closes the review, releases the hold, records the outcome		
5	(H) Warm handoff: specialist joins with transcript, verification result, sentiment flag		



J ELEVENLABS COMPONENTS		Tick + 60 words · Agent design · 25%	
Tick only what your use case genuinely needs, then justify the two least obvious choices. Ticking everything scores lower than ticking four with reasons.			
☐ Agents Platform	☐ Agent Workflows	☐ Sub-agents	☒ Eleven v3 TTS
☐ Voice Design	☒ Scribe v2 STT	☐ Knowledge base + RAG	☐ Server / client tools
☐ MCP servers	☐ Telephony (Twilio / SIP)	☐ Batch calling	☐ Agent Testing
☐ Post-call webhooks	☐ WhatsApp	☐ Web / mobile SDKs	☐ Bring-your-own LLM
Why these two: STT has to hear 'not mine' in six languages (scribe_v2); TTS has to say the disclosure back in the customer's own language.			

K GUARDRAILS		20 words per row · Guardrails · 20%	
Every row must name a mechanism in your design, not an intention. "We will be careful" scores zero.			
REQUIREMENT	HOW YOUR DESIGN ENFORCES IT		
Opening disclosure	auditAgentReply scans every reply (compliance/policy.ts:65); a missing turn-1 disclosure is injected server-side before speech		
Consent to be called	enroll returns 422 without consentRecordId (enroll/route.ts:132); requireOutboundConsent gates every outbound call		
Verification without secrets	credential deny-list (policy.ts:53) refuses PIN, OTP, CVV, PAN and substitutes a safe refusal - server-side, no toggle		
Human approval point	planFor() (interventions/route.ts:71) hands off to fraud_specialist at every tier; the agent may only hold temporarily		
Opt-out path	optedOut on the customer (enroll/route.ts:110) is checked before delivery; re-enroll is deliberate re-consent		
Escalation trigger	sentiment.ts flags distress words, third-party pressure, repetition and forces handoff whatever the classifier said		

04 The Architecture

What you are actually building, and what it connects to.

L TECHNICAL ARCHITECTURE		Diagram · Agent design · 25%	
Draw your system across the three zones. Every arrow must be labelled with what flows across it and in which direction. Mark each boundary where personal data crosses with a filled dot.			
Caller and channel <i>Phone, web, WhatsApp, inbound or outbound</i>	C1. Browser demo /api/tts + /api/asr speechSynthesis fallback		
	C2. Customer phone (PSTN) Twilio <Play> ElevenLabs MP3 <Say> voice on each turn risk signal down: ref, amount, merchant, score. no card data	C3. Bank fraud engine POST /v1/interventions, HMAC-signed enroll + opt-out via /v1/enroll customer speech down; phone number already held by the bank (PD) audited reply up, spoken in the customer's language	
ElevenLabs platform <i>Agent, voice, knowledge, tools, testing</i>	2. Eleven v3 TTS (out) eleven_multilingual_v2 today eleven_flash_v2_5 for Swahili	3. Compliance guardrails policy.ts auditAgentReply credential deny-list	
	1. Scribe v2 STT (in) /api/asr, scribe_v2 Deepgram / z-ai only as substitutes	4. Groq rephrase layer optional, llm.ts still audited after	
	0. Agent logic + intent router /api/agent, /api/twilio/turn deterministic intent + sentiment temporary freeze request down, ref only	5. Audit chain (SHA-256) audit-chain.ts transcript redacted before store hash-chained rows down: redacted transcript + every action (PD) specialist confirm / reverse up: only a human releases a permanent block	
Institution systems <i>Core systems, data stores, human queues</i>	I1. Fraud specialist queue HUMAN APPROVAL GATE confirm or reverse the hold	I2. Core banking card-freeze API temporary hold only, never permanent	
	I3. CRM / customer records optedOut + consentRecordId	I4. Postgres via Prisma cases, intervention + audit rows	
Your diagram must show:			
☒ Every ElevenLabs component you ticked in box J		☒ Each institution system by name	
☒ The direction of flow on every arrow		☒ Where personal data crosses a boundary	

☒ The human approval gate from box K☒ What happens when a dependency is down**Box L continued - the human approval gate named in box K, and what happens when a dependency is down****Dependency failure modes - what degrades, what never happens**

05 The Case

What success looks like, what could go wrong, and why this team.

M SUCCESS METRICS	Max 3 KPIs · Commercial · 20%																
<i>Each baseline must match box D exactly. Three sharp metrics beat eight vague ones.</i>																	
<table><thead><tr><th>KPI</th><th>BASELINE (FROM D)</th><th>TARGET</th><th>HOW IT IS MEASURED</th></tr></thead><tbody><tr><td>Time from fraud signal to customer contact</td><td>38 min</td><td>< 60 sec</td><td>receivedAt to Twilio call start_time, per case</td></tr><tr><td>Languages live on fraud calls</td><td>2 (EN/AR)</td><td>6 langs</td><td>languages with full script + voice + intent keywords</td></tr><tr><td>Unauthorized-transaction write-off rate</td><td>1.0%</td><td>0.3%</td><td>monthly chargeback write-offs / card volume</td></tr></tbody></table>	KPI	BASELINE (FROM D)	TARGET	HOW IT IS MEASURED	Time from fraud signal to customer contact	38 min	< 60 sec	receivedAt to Twilio call start_time, per case	Languages live on fraud calls	2 (EN/AR)	6 langs	languages with full script + voice + intent keywords	Unauthorized-transaction write-off rate	1.0%	0.3%	monthly chargeback write-offs / card volume	
KPI	BASELINE (FROM D)	TARGET	HOW IT IS MEASURED														
Time from fraud signal to customer contact	38 min	< 60 sec	receivedAt to Twilio call start_time, per case														
Languages live on fraud calls	2 (EN/AR)	6 langs	languages with full script + voice + intent keywords														
Unauthorized-transaction write-off rate	1.0%	0.3%	monthly chargeback write-offs / card volume														
N RISKS	25 words per row · Guardrails · 20%																
<i>Deployment, conduct or regulatory. Name the risk you are most likely to be challenged on — not the safest one.</i>																	
<table><thead><tr><th>RISK</th><th>HOW YOU HANDLE IT</th></tr></thead><tbody><tr><td>1 AI caller reads as spam, customer hangs up</td><td>Disclosure names the bank and the recording purpose in sentence one; opt-out is one word away</td></tr><tr><td>2 Fraudster spoofs the bank and calls first</td><td>The agent never asks for credentials; a spoof always does; the audit chain proves the real call</td></tr><tr><td>3 CBUAE challenges an AI's authority to freeze</td><td>Freeze is temporary and reversible, irreversible steps need the specialist, every action is hash-chained</td></tr></tbody></table>	RISK	HOW YOU HANDLE IT	1 AI caller reads as spam, customer hangs up	Disclosure names the bank and the recording purpose in sentence one; opt-out is one word away	2 Fraudster spoofs the bank and calls first	The agent never asks for credentials; a spoof always does; the audit chain proves the real call	3 CBUAE challenges an AI's authority to freeze	Freeze is temporary and reversible, irreversible steps need the specialist, every action is hash-chained									
RISK	HOW YOU HANDLE IT																
1 AI caller reads as spam, customer hangs up	Disclosure names the bank and the recording purpose in sentence one; opt-out is one word away																
2 Fraudster spoofs the bank and calls first	The agent never asks for credentials; a spoof always does; the audit chain proves the real call																
3 CBUAE challenges an AI's authority to freeze	Freeze is temporary and reversible, irreversible steps need the specialist, every action is hash-chained																
O WHAT WILL BE WORKING BY 14 OCTOBER	60 words max · Commercial · 20%																
<i>Two weeks on the platform. What genuinely runs end to end, and what stays a mock? Checked at Stage 2.</i> 18 of 18 evidence-pack checks pass live against the running app: HMAC signal intake, idempotent replay, six-language intent routing, disclosure and credential guardrails, rate limiting, ElevenLabs TTS, Scribe v2 STT and hash-chained audit read-back. Twilio runs in api-key mode. Unproven: one real outbound dial to a handset.																	
P TEAM	Structured · Team · 10%																
<i>Who builds this, and what each person has shipped before.</i>																	
<table><thead><tr><th>NAME</th><th>ROLE ON THIS BUILD</th><th>SHIPPED PREVIOUSLY (LINK)</th></tr></thead><tbody><tr><td>Aaron Otema</td><td>Solo founder: architecture, agent, compliance layer, telephony</td><td>This repo: github.com/BROCKUGANDA/SecureVoice-AI</td></tr><tr><td></td><td></td><td></td></tr><tr><td></td><td></td><td></td></tr></tbody></table>	NAME	ROLE ON THIS BUILD	SHIPPED PREVIOUSLY (LINK)	Aaron Otema	Solo founder: architecture, agent, compliance layer, telephony	This repo: github.com/BROCKUGANDA/SecureVoice-AI											
NAME	ROLE ON THIS BUILD	SHIPPED PREVIOUSLY (LINK)															
Aaron Otema	Solo founder: architecture, agent, compliance layer, telephony	This repo: github.com/BROCKUGANDA/SecureVoice-AI															
Q PROOF OF BUILD	Two links · Team · 10%																
<i>One link to something your team has deployed. One to a 60-second recording of a team member walking through your box L diagram. Both required.</i> Deployed build: github.com/BROCKUGANDA/SecureVoice-AI 60-second walkthrough of box L: https://youtu.be/DFAHZzaAxyz																	

How this canvas is scored

Criterion	Weight	Evidence comes from
Problem fit and depth of institutional understanding	25%	Boxes B, C, D, F, G, H
Agent design and use of the ElevenLabs stack	25%	Boxes I, J, L
Guardrail and compliance design	20%	Boxes K, N
Commercial viability and readiness to build	20%	Boxes E, M, O
Team	10%	Boxes P, Q